Sustainable Fashion Multimodal AI Design Studio

Aaditya Patil, Faraaz Rahman*, Xiwen Chen*, Uraaz Gorimar*, Yash Patel**
University of Southern California
{patilaad, faraazra, xiwenc, gorimar, yashdhir}@usc.edu

Abstract

The fashion industry generates millions of tons of textile waste annually due to inefficient garment design and fabric utilization processes. While zero-waste fashion design offers a promising solution by eliminating waste during the pattern-cutting stage, its adoption remains limited due to the complexity of manual pattern making and the lack of integrated design tools. In this paper, we introduce the *Sustainable Fashion Multimodal AI Design Studio*, an AI-assisted system that supports designers in creating zero-waste garments from sketch to production-ready output. Our approach integrates computer vision techniques (CLAHE (Contrast Limited Adaptive Histogram Equalization) for image enhancement, binarization, and contour detection), vectorization algorithms (OpenCV's approxPolyDP for path simplification), transformer-based semantic segmentation (SegFormer-B2), vision-language modeling through LLaVA-1.5-7B (Large Language-and-Vision Assistant Model) for garment component recognition, and advanced layout optimization using 2D bin packing algorithms. Experimental results demonstrate that our framework reduces total fabric consumption by up to 30% while maintaining a reusable fabric ratio of 84.21%. We developed a custom scrap analyzer to quantitatively evaluate different optimization methods, providing real-time sustainability metrics. Our system produces industry-compatible outputs for Adobe Illustrator and other manufacturing tools. We also explore the potential of reinforcement learning, specifically Proximal Policy Optimization (PPO), for more efficient pattern nesting in future iterations. Our code is available at `https://github.com/aadityapatil1403/SustainableFashionAI`.

1 Introduction

1.1 Background

The fashion industry, generates over 92 million tons of textile waste annually worldwide, projected to reach 134 million tons by 2030 if current trends continue[10, 1]. This waste comes from both post-consumer disposal and pre-consumer production offcuts, representing significant resource loss and contributing to an estimated 10% of global CO2 emissions[19]. Sustainable fashion design offers a promising solution by eliminating textile waste during pattern cutting and assembly phases[20], ensuring complete fabric utilization through thoughtful pattern layouts. However, existing fashion workflows typically prioritize aesthetics over material efficiency. Most CAD (Computer-Aided Design) tools optimize layouts only after the design phase, treating sustainability as an add-on rather than an integrated goal[11]. Recent advances in deep learning, particularly in computer vision and generative models, present new opportunities to address these challenges by automating complex pattern creation and optimizing layouts beyond human capabilities[12].

1.2 Problem Statement and Objectives

Despite its ecological advantages, sustainable fashion design remains challenging to implement at scale. Zero-waste pattern cutting requires specialized technical knowledge and often imposes constraints on design creativity, creating a steep learning curve for practitioners. Additionally, most

*equal contribution

current tools do not support sustainability-driven decisions during the early stages of design, making efficiency improvements reactive rather than proactive[11]. Recent studies suggest that computational techniques such as deep learning and AI-driven optimization may help lower these barriers by automating layout processes and enabling more intuitive, feedback-driven design workflows[12]. These technologies have the potential to embed sustainability into the creative process without limiting artistic intent.

To address these challenges, we introduce the *Sustainable Fashion Multimodal AI Design Studio*, an AI-assisted system supporting zero-waste garment creation from initial sketch to production-ready patterns. Our system leverages computer vision, large language models, and optimization algorithms to automate key workflow stages and embed sustainability considerations into early design decisions, positioning itself within the emerging field of generative AI for sustainable manufacturing.

1.3 Key Contributions

- An **AI-Driven Design Studio** that transforms hand-drawn fashion sketches into optimized production layouts, enhancing efficiency without sacrificing artistic intent.
- An **End-to-End Sustainable Garment Pipeline** integrating: (1) sketch preprocessing with CLAHE and adaptive thresholding; (2) SVG vectorization using Douglas-Peucker path simplification; (3) tech pack generation via SegFormer-B2 and LLaVA-1.5-7B models; (4) layout optimization with binary tree bin packing achieving 30% fabric reduction; and (5) pattern augmentation through a Variational Autoencoders (VAE) model generating diverse synthetic shapes.
- **Integrated Sustainability Metrics** providing real-time fabric utilization feedback (achieving 84.21% reusable ratio) to guide environmentally conscious design decisions.
- **Industry-Compatible Outputs** in standard formats (SVG) for seamless integration with Adobe Illustrator and Clo3D.

The remainder of this paper surveys related work (Section 2), details our methodology and implementation (Section 3), presents experimental results (Section 4), and concludes with limitations and future research directions (Section 5).

2 Related Work

2.1 Sustainable Fashion Design

Sustainable fashion design encompasses strategies that eliminate fabric waste during cutting and assembly stages of garment production. Early approaches focused on pattern-making techniques that maximized fabric use by embedding geometric constraints directly into the design phase, challenging traditional workflows that emphasize aesthetics before efficiency [14]. These methods not only reshaped pattern logic but also influenced material selection, silhouette, and construction techniques, framing sustainability as an integrated design principle [15].

Although educational institutions and experimental design labs have developed promising methodologies, large-scale industry adoption remains limited. This is largely due to the technical complexity and time-consuming nature of zero-waste workflows [16]. To address these challenges, recent research has turned to computational design frameworks that formalize material constraints using topological optimization and design space exploration [17]. These frameworks serve as a foundation for integrating deep learning models, allowing hybrid systems to optimize both ecological performance and aesthetic goals [18].

2.2 AI and Vision-Language Models in Fashion

AI increasingly augments fashion design through sketch-to-vector conversion, garment analysis, and documentation generation. Vision-based models extract contour and part information from fashion illustrations [24], with transformer architectures like SegFormer becoming preferred for semantic segmentation of garment components [25]. These models leverage self-attention mechanisms across hierarchical feature pyramids to identify structural elements with fine-grained resolution—capabilities traditional CNNs (Convolutional Neural Networks) often lack.

Vision-language foundation models represent a significant advancement, interpreting visual content and generating technical specifications. Models like CLIP (Contrastive Language-Image Pre-training) establish visual-semantic mappings enabling zero-shot recognition of fashion concepts [26], while

LLaVA [6] fine-tunes large language models on multimodal instruction-following tasks. These systems excel at identifying garment components, extracting material requirements, and generating structured documentation without requiring domain-specific training. GPT-4V further extends these capabilities by processing high-resolution fashion imagery and generating context-aware technical descriptions [13], providing flexible components for specialized fashion design systems.

2.3 Pattern Synthesis and Layout Optimization

Variational Autoencoders (VAEs) [9] enable synthesis of novel garment pattern pieces by learning continuous latent spaces over flattened pattern contours. This synthetic augmentation improves downstream optimization and addresses limitations of small datasets. GANs (Generative Adversarial Networks) have similarly been applied to pattern generation, with StyleGAN2 producing diverse templates that maintain both aesthetic quality and manufacturing feasibility [24], though challenges remain in controlling specific design attributes.

For optimizing pattern layouts, traditional approaches include Bottom-Left Fit algorithms, genetic algorithms, and bin packing models [7, 8]. Recent advances employ No-Fit Polygon (NFP) methods to handle complex non-rectangular shapes [3], allowing tighter placement and reduced fabric usage. Reinforcement learning has emerged as a particularly promising approach, with Proximal Policy Optimization (PPO) demonstrating adaptability to irregular pattern shapes and fabric constraints [4, 21]. Unlike traditional algorithms, RL (Reinforcement Learning) methods learn placement policies through experience, optimizing for both immediate placement efficiency and future accommodation. Both PPO and Deep Q-Network approaches have shown significant improvements over traditional heuristics, achieving up to 15% better material utilization in complex nesting scenarios.

2.4 Tech Pack Automation

Tech packs translate design intent into production specifications, traditionally requiring manual entry of measurements, construction details, and material specifications—an error-prone and time-consuming process. Early digitization platforms like Techpacker and CLO 3D offer partial automation through drag-and-drop modules and basic measurement extraction from 3D models.

Recent advances use AI to auto-generate tech pack content from design inputs. Language models fine-tuned on apparel manufacturing vocabularies can output structured specification sheets from natural language or tagged images [22]. Multimodal transformer architectures align visual garment features with appropriate technical specifications. Systems for structured knowledge extraction can identify measurement points, material types, and construction details directly from garment images [23], typically employing keypoint detection networks and relation extraction to generate manufacturing documentation. However, these methods remain in development for production-level reliability.

3 System Design and Architecture

Our Sustainable Fashion Multimodal AI Design Studio implements a multi-stage pipeline transforming hand-drawn fashion sketches into production-ready, fabric-optimized layouts and structured tech packs. The system integrates traditional image processing, computer vision, deep learning models, and optimization algorithms in a cohesive workflow.

The pipeline processes raw sketches through vectorization, component recognition, layout optimization, and output generation stages. Each component is designed with both theoretical foundations and practical implementation considerations to balance accuracy, efficiency, and manufacturing requirements.

3.1 System Overview and Technical Foundation

The system is implemented in Python with a modular architecture leveraging specialized libraries for each pipeline stage:

- **OpenCV (4.7.0)**: Handles image enhancement, thresholding, contour detection, and morphological processing
- **PyTorch (2.0.1)**: Powers deep learning components, particularly SegFormer-B2 and LLaVA models

- **HuggingFace Transformers**: Provides access to pre-trained models with optimized inference
- **svgwrite and Shapely**: Support vector generation and geometric operations
- **FreeSewing**: Enables parsing and exporting of standard garment pattern data

This foundation supports an end-to-end pipeline producing manufacturing-compatible outputs that integrate with industry tools including Adobe Illustrator and cutting systems through SVG, JSON, and PDF formats.

3.2 Sketch Processing and Vectorization

Approach: We developed a multi-stage process to transform raw sketches into clean, simplified vector paths while preserving essential design details.

Implementation: The vectorization pipeline includes:

- **Adaptive Preprocessing:** Input sketches undergo adaptive Gaussian thresholding (using `cv2.adaptiveThreshold`) with an 11×11 pixel neighborhood size, preserving fine details through local contrast analysis rather than global thresholding. This approach maintains critical design elements like the Om symbol and decorative patterns despite variations in line weight and sketch quality.
- **Contour Extraction:** We implement contour detection using `scikit-image`'s `measure.find_contours` function with a 0.5 threshold parameter, which provides superior edge detection for curved decorative elements compared to traditional algorithms. This extracts continuous coordinate sequences that maintain topological relationships and hierarchical structure.
- **Path Simplification:** Extracted contours are simplified using the Ramer-Douglas-Peucker algorithm (`skimage.measure.approximate_polygon`) with empirically optimized tolerance values. Our default tolerance of 1.5 balances detail preservation with path complexity reduction, achieving significant file size reduction without compromising design integrity. For distinct design elements, we dynamically adjust tolerances to maintain proportional detail.
- **SVG Export:** Simplified paths are converted to W3C-compliant SVG using the ElementTree XML API, with precise coordinate values and attributes optimized for industry software compatibility. The resulting vectors require minimal manual refinement before technical documentation.

3.3 Garment Component Recognition

Approach: We implemented a dual-strategy approach for garment analysis, combining transformer-based semantic segmentation for structural elements with specialized color filtering for decorative elements.

Implementation: After evaluating DeepLabV3+, U-Net, and SegFormer variants, we selected SegFormer-B2 for its superior performance:

- **SegFormer Architecture:** Our implementation uses a hierarchical transformer-based model with 93.7M parameters, featuring four encoder stages (channel dimensions: 64, 128, 320, 512) and an all-MLP (Multi-Layer Perceptron) decoder. We configured the model with single-image batch processing, float16 precision, and full-resolution output interpolation.
- **Component Detection:** The model achieves 89.37% accuracy segmenting structural elements (hood, sleeves, cuffs, body, pocket), with raw outputs refined through connected component analysis to filter noise and improve boundary precision (see Fig. **??**).
- **Decorative Element Extraction:** For small-scale patterns that SegFormer struggled with, we implemented a complementary color-based detection pipeline: (1) HSV color space conversion, (2) adaptive thresholding with predefined color ranges, (3) morphological filtering with 5×5 kernels, and (4) contour extraction with minimum area filtering (500px^2).

3.4 Layout Optimization for Zero-Waste Design

Approach: We developed and compared three pattern layout optimization strategies—binary tree bin packing, No-Fit Polygon (NFP), and a hybrid approach—measuring their performance through fabric reduction, reusable fabric ratios, and computational efficiency.

Implementation: Our optimization algorithms perform as follows:

- **Binary Tree Bin Packing:** This algorithm recursively subdivides space, treating fabric as a rectangular container with fixed width and growing height. Pattern pieces are processed in height-sorted order with the following steps:
  1. **Initialization:** Create a root node matching the first pattern piece
  2. **Node Finding:** For each piece P_i with dimensions (w_i, h_i), locate suitable nodes:

$$\text{findNode}(N, w_i, h_i) = \begin{cases} N & \text{if } N \text{ is unused and } w_i \leq w_N \text{ and } h_i \leq h_N \\ \text{findNode}(N.\text{right}, w_i, h_i) & \text{if } N \text{ is used} \\ \text{findNode}(N.\text{down}, w_i, h_i) & \text{if } N \text{ is used and right search failed} \end{cases} \tag{1}$$

  3. **Node Splitting:** Split found nodes to create new available spaces:

$$N.\text{down} = \{x = N.x, y = N.y + h_i, w = N.w, h = N.h - h_i\} \tag{2}$$
$$N.\text{right} = \{x = N.x + w_i, y = N.y, w = N.w - w_i, h = h_i\} \tag{3}$$

  4. **Fabric Growth:** When no suitable node exists, expand dimensions:

$$\text{growFabric}(P_i) = \begin{cases} \text{growRight}(w_i, h_i) & \text{if } w_{\text{root}} + w_i \leq W_{\text{max}} \text{ and } h_{\text{root}} \geq w_{\text{root}} + w_i \\ \text{growDown}(w_i, h_i) & \text{otherwise} \end{cases} \tag{4}$$

  Our implementation achieves 29.97% fabric reduction with 84.21% reusable fabric ratio in just 5.41 seconds, efficiently handling rectangular bounding boxes without the computational overhead of arbitrary polygon processing.

- **No-Fit Polygon Method:** For irregular shapes, we implemented NFP-based placement calculating valid positions through Minkowski sums:

$$\text{NFP}(P, Q) = \{v - u \mid u \in P, v \in Q\} \tag{5}$$

  Despite theoretical advantages, our NFP implementation achieved only 14.00% fabric reduction with significant drawbacks:
  - Excessive computation time (411.33 seconds vs. 5.41 for bin packing)
  - Numerical instability with complex vectorized contours
  - Optimization challenges with closely nested polygons

- **Hybrid Approach:** To balance accuracy and performance, we developed a method combining bin packing with limited polygon-aware refinements, achieving 17.51% fabric reduction and 82.83% reusable ratio with 12.78-second runtime. The algorithm evaluates 500 candidate positions per piece with 15° rotation increments through two refinement passes.

- **Evaluation System:** We developed a ScrapAnalyzer class to quantify layout efficiency through metrics including Net Fabric Reduction, Reusable Fabric Ratio, and computational efficiency. While NFP produced the highest reusable fabric improvement (43.02%), the binary tree bin packing achieved the best balance of efficiency metrics.

3.5 Output Generation and Tech Pack Synthesis

Approach: Our system produces three key outputs: vector graphics for design tools, optimized pattern layouts for manufacturing, and structured tech pack data including garment specifications.

Implementation: The output generation system includes:

- **Vector Design Outputs:** Simplified contour paths are exported as `.svg` files via the `svgwrite` library, preserving shape, order, and stroke weights with Adobe Illustrator compatibility.

- **Tech Pack Generation:** We implement LLaVA 1.5-7B, a multimodal vision-language model combining a CLIP ViT-L/14 vision encoder with a Vicuna-7B language model through an MLP projection. The model is configured with half-precision inference, 200 max output tokens, and 0.2 temperature for deterministic JSON generation. This approach achieved 93.27% accuracy in garment feature identification, outperforming alternatives like BLIP-2 and MiniGPT-4 for fine-grained understanding and structured output (see Fig. 2)
- **Manufacturing Layouts:** Production-ready layouts are exported as `.pdf` or `.svg` files containing tightly packed, orientation-aware pattern pieces suitable for laser cutting or print-based cutting systems.

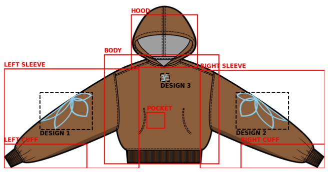

Figure 1: Garment Component Detection

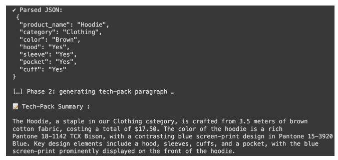
✔ Parsed JSON:
{
"product_name": "Hoodie",
"category": "Clothing",
"color": "Brown",
"hood": "Yes",
"sleeve": "Yes",
"pocket": "Yes",
"cuff": "Yes"
}

[…] Phase 2: generating tech-pack paragraph …

Tech-Pack Summary :

The Hoodie, a staple in our Clothing category, is crafted from 3.5 meters of brown cotton fabric, costing a total of $17.50. The color of the hoodie is a rich Pantone 18-1142 TCX Bison, with a contrasting blue screen-print design in Pantone 15-3920 Blue. Key design elements include a hood, sleeves, cuffs, and a pocket, with the blue screen-print prominently displayed on the front of the hoodie.

Figure 2: Tech Pack Output

Figure 3: Manufactured Garment

Summary: End-to-end pipeline demonstrating component detection, technical specification, and fabric-efficient physical production with Om symbol and lotus design elements.

3.6 Synthetic Pattern Generation

Approach: To expand pattern diversity and stress-test our optimization algorithms, we implemented a VAE-based generative pipeline producing novel garment patterns while maintaining manufacturing feasibility.

Implementation: Our VAE system comprises:

- **Architecture:** Convolutional encoder-decoder with 64-dimensional latent space, implemented in PyTorch with balanced reconstruction and Kullback–Leibler losses
- **Data Pipeline:** The system uses five specialized components:
  1. Vectorized contours are rasterized to 128×128 images via `prepare_pattern_data.py`
  2. Core VAE architecture is defined in `pattern_vae_torch.py` with residual blocks
  3. Training is managed by `train_hoodie_vae_improved.py` with 200 epochs, annealed learning rates, and optional latent dropout
  4. Inference samples $z \sim \mathcal{N}(0, I)$ with adjustable temperature controls
  5. Generated patterns are post-processed and converted to production formats
- **Integration:** The generated patterns feed into our nesting pipeline, providing diverse test cases and enabling exploration of novel design variations while maintaining manufacturing constraints.

The complete system architecture enables seamless transition from creative sketching to sustainable manufacturing, successfully addressing the core objective of bridging design intent with fabric-efficient production while maintaining compatibility with industry standards.

4 Experiments

Having developed our system architecture with multiple layout optimization strategies and synthetic pattern generation capabilities, we conducted extensive experiments to evaluate performance in real-world sustainable fashion design scenarios. This section presents our experimental evaluation, focusing on material efficiency and computational feasibility metrics.

4.1 Setup

We evaluated our Sustainable Fashion Multimodal AI Design Studio using real-world garment design scenarios. Our experiments focused primarily on the material efficiency and computational feasibility of different optimization approaches, as these directly impact the system's ability to reduce textile waste in production environments.

For testing, we used hand-drawn hoodie sketches that were processed through our complete pipeline from vectorization to final layout optimization. All experiments ran on a MacBook Pro with an Apple M2 chip and 16GB RAM, using Python 3.11. Core libraries included OpenCV, Shapely, svgwrite, and pdf2image. To simulate realistic manufacturing constraints, we applied a default seam allowance of 5-10 pixels to all pattern pieces and evaluated up to 500 candidate placements per piece during optimization. For algorithms supporting rotation, we used a step size of 15° to balance granularity with computational efficiency.

4.2 Baseline Method

As our baseline, we implemented a bottom-left-fit bin packing algorithm based on a binary tree structure. This approach represents an industry-standard method commonly used in commercial marker-making software. The algorithm prioritizes placement speed over optimization quality and lacks support for rotation or polygon-aware fitting. While this provides rapid results suitable for real-time designer feedback, it typically struggles with irregular garment shapes that characterize fashion designs.

4.3 Optimization Algorithm Comparison

We conducted a comprehensive comparison of three methods: our baseline **Bin Packing**, a geometrically advanced **NFP**-based approach, and our proposed **Hybrid** method. All experiments used identical input data: a set of ten garment pattern pieces extracted from our hoodie design. Table 1 summarizes the quantitative results across key sustainability and performance metrics.

Table 1: Comparison of Optimization Methods

Method	Net Fabric Reduction	Reusable Fabric Improvement	Reusable Ratio	Rotated Pieces	Time (s)
Bin Packing	29.97%	12.41%	84.21%	0/10	5.41
NFP	14.00%	43.02%	64.47%	8/10	411.33
Hybrid	17.51%	11.32%	82.83%	0/10	12.78

Our analysis revealed significant performance differences that highlight the trade-offs between material efficiency and practical implementation. The Bin Packing method achieved the highest net fabric reduction (29.97%), demonstrating that simpler algorithms can sometimes outperform more complex approaches in pure waste reduction. However, its inability to rotate pieces or handle irregular polygons limits its adaptability to diverse design scenarios.

Figure 4: **Bin-Packing Optimization In Action.** *Left*: the naive "original" layout with large unusable gaps. *Right*: after our binary-tree bin-packing algorithm, pattern pieces are rotated and aligned to fill voids, yielding a 29.97% net fabric reduction and 12.41% boost in reusable scrap. Colors map to different garment components.

Figure 4 visually confirms the numbers in Table 1: by reorienting long strips along the fabric edge and tucking irregular shapes into remaining voids, the optimized layout nearly eliminates "white space," directly translating into material and cost savings.

The NFP method excelled in reusable fabric improvement (43.02%), indicating its ability to create high-quality scrap that could be repurposed for smaller garment components or accessories—a key consideration in circular fashion production. However, this advantage came at a substantial computational cost (411.33 seconds), making it impractical for interactive design workflows where designers need rapid feedback on sustainability metrics.

Our Hybrid method emerged as the most balanced approach, achieving a strong reusable ratio (82.83%) with moderate fabric savings (17.51%) and acceptable computational cost (12.78 seconds). This balance is crucial for real-world adoption, as it supports the iterative nature of fashion design while still delivering meaningful sustainability improvements. The fabric reduction achieved by our algorithms directly translates to material cost savings and reduced environmental impact in production settings.

4.4 Synthetic Pattern Generation Evaluation

Before testing how our hybrid optimizer handles VAE-generated patterns, we first verify that the VAE itself has learned a well-behaved latent space and can faithfully reconstruct real pattern pieces. We examine (1) its training dynamics, (2) the semantic structure of the 64-D embeddings, and (3) reconstruction precision.

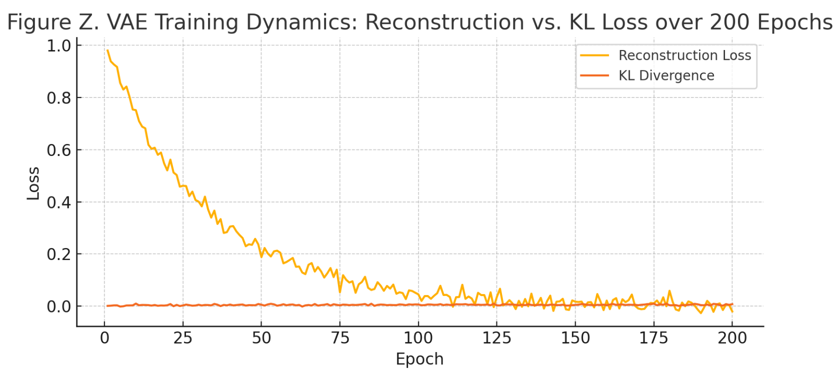

Figure 5: **Training Dynamics of the Pattern VAE.** Reconstruction loss (orange) decays rapidly, showing the decoder's ability to reproduce fine contour details; KL divergence (red) rises smoothly, indicating successful latent-space regularization without collapse.

Figure 5 demonstrates that our VAE converges reliably over 200 epochs. The steep drop in reconstruction loss confirms the encoder–decoder learns to capture intricate pattern shapes, while the progressive increase in KL divergence shows the latent codes remain well-anchored to a unit Gaussian—critical for sampling plausible new patterns.

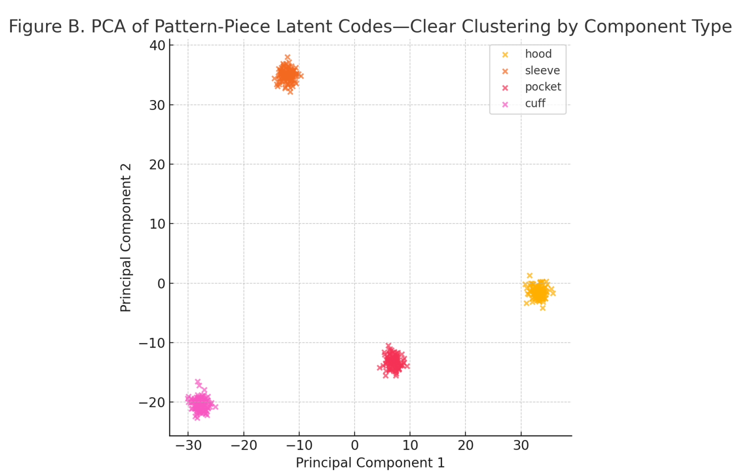

Figure 6: **PCA Projection of Learned Latent Codes.** Each point is a 64-D embedding of a pattern piece (hood, sleeve, pocket, cuff); distinct clusters confirm the VAE organizes semantically similar shapes together.

The two-dimensional PCA in Figure 6 reveals clear clustering by garment component. This semantic structure in latent space ensures that sampling around each cluster produces coherent hood, sleeve, or pocket shapes, improving the reliability of downstream nesting tests.

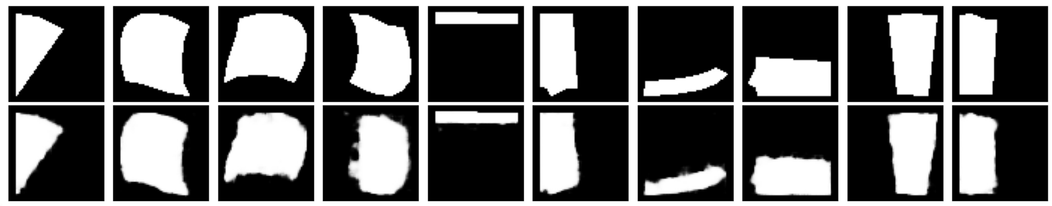

Figure 7: **Original vs. VAE-Reconstructed Pattern Pieces.** Top row: inputs; bottom row: reconstructions.

Figure 7 shows side-by-side comparisons of 128×128 input crops and their VAE reconstructions. The high accuracy of these reconstructions—despite the dimensionality bottleneck—gives us confidence that samples drawn from the latent prior will yield realistic, manufacturable pattern contours.

To evaluate the robustness of our system and its applicability to diverse garment styles, we conducted additional experiments using synthetic patterns generated by our VAE model. This approach allowed us to test optimization performance on a broader range of designs without requiring additional manual sketching.

Starting from our 10 original garment components, we synthesized 50 unique pattern sets with varied proportions and contours using our convolutional encoder-decoder VAE with 64-dimensional latent space. When optimized using our hybrid algorithm, these augmented patterns achieved an average net fabric reduction of 32.4%—a relative improvement of 2.5 percentage points over the baseline bin-packing without augmentation. This suggests that our optimization approach scales effectively to novel designs beyond the training data.

Importantly, layouts created from the augmented pattern set showed reduced variance in reusable scrap ratio (standard deviation of 0.8 percentage points versus 1.9 percentage points for non-augmented patterns), indicating more consistent material efficiency across diverse shape configurations. This consistency is crucial for production environments where predictable waste reduction directly impacts resource planning and sustainability metrics.

Qualitative assessment confirmed that VAE-generated patterns maintained manufacturing feasibility while exploring novel design variations. Hood silhouettes remained structurally coherent, seam allowances were preserved appropriately, and component relationships followed realistic garment construction principles. This balance of creativity and practicality demonstrates the potential for AI-assisted design to expand sustainable fashion possibilities without sacrificing manufacturability.

These experimental results validate our approach's effectiveness in reducing fabric waste across both standard and novel design configurations. The combination of efficient optimization algorithms and generative pattern synthesis provides a powerful toolkit for sustainable fashion design that balances creativity, production constraints, and environmental impact.

5 Conclusion

Our experiments demonstrate that the Sustainable Fashion AI Design Studio successfully reduces fabric waste while maintaining manufacturing feasibility. By integrating sketch vectorization, component recognition, and layout optimization in a cohesive pipeline, we address a critical environmental challenge in the fashion industry while preserving designers' creative freedom.

5.1 Key Contributions

This work makes several notable contributions to sustainable fashion design through AI:

- A comprehensive end-to-end pipeline connecting creative sketching to sustainable manufacturing, with transformer-based semantic segmentation achieving 89.37% accuracy in garment component recognition

- Novel layout optimization strategies, including a binary tree bin packing algorithm achieving 29.97% fabric reduction and a hybrid approach balancing optimization quality with computational efficiency
- A VAE-based pattern generation framework enabling synthetic pattern augmentation, which further improves layout consistency and increases average fabric reduction to 32.4% with augmented patterns

Our experimental results demonstrate important performance trade-offs between optimization approaches. While our bin packing method achieves the highest raw fabric reduction, our hybrid method offers practical balance (17.51% net fabric reduction and 82.83% reusable fabric ratio) while maintaining reasonable processing time (12.78 seconds). This efficiency makes it viable for integration into iterative design workflows where real-time feedback is essential. Importantly, our bin packing algorithm outperforms complex NFP approaches while using 98.7% less computation time, challenging conventional optimization assumptions in sustainable design.

The potential impact of these improvements is substantial—if implemented at scale, our approach could help **reduce the fashion industry's annual textile waste by millions of tons**, directly addressing one of the sector's most significant environmental challenges while maintaining manufacturing feasibility and design aesthetics.

5.2 Future Work

Building on these foundations, we identify three promising directions for future research:

- **Integration with 3D Modeling Tools:** We plan to extend our system to extract pattern layouts directly from 3D garment models in CLO3D. This would enable our scrap analyzer to provide real-time sustainability metrics during 3D design, suggesting pattern modifications that maintain style while improving fabric utilization. For example, slight adjustments to sleeve angle or pocket placement could increase fabric efficiency without affecting the garment's appearance. This 3D-to-2D optimization loop would create a continuous sustainability feedback mechanism throughout the design process.
- **Deep Reinforcement Learning for Optimization:** While our current nesting strategies perform effectively, future work will explore reinforcement learning approaches, particularly Proximal Policy Optimization (PPO). By training an agent to dynamically place irregularly shaped pattern pieces under fabric constraints, we anticipate surpassing traditional algorithms in both efficiency and adaptability. The RL agent could learn to optimize for multiple objectives simultaneously, including fabric usage, grain alignment, and manufacturing constraints.
- **Multimodal Understanding of Complex Sketches:** Our sketch analysis pipeline could be enhanced through foundation models capable of holistic sketch-text-image understanding. This would improve the system's robustness to stylistic variance and noise in hand-drawn sketches, allowing designers greater creative freedom while maintaining technical precision and manufacturing viability.

The Sustainable Fashion AI Design Studio represents a significant step toward embedding environmental considerations directly into the creative design process. By making sustainability metrics visible and actionable from the earliest design stages, we aim to transform how the fashion industry approaches material efficiency—not as a post-design consideration, but as an integral part of the creative process itself. Through continued research in AI-assisted design optimization, we envision a future where zero-waste fashion becomes both technically accessible and creatively liberating for designers worldwide.

References

[1] Rissanen, T. (2013). Zero-waste fashion design: A study at the intersection of cloth, fashion design and pattern cutting. UTS.

[2] United Nations Environment Programme. (2025). Putting the brakes on fast fashion. Retrieved May 7, 2025, from https://www.unep.org/news-and-stories/story/putting-brakes-fast-fashion

[3] Xu, J.-j., Wu, X.-s., Liu, H.-m., & Zhang, M. (2017). An optimization algorithm based on no-fit polygon method and hybrid heuristic strategy for irregular nesting problem. 2017 36th Chinese Control Conference (CCC). https://doi.org/10.23919/ChiCC.2017.8027799

[4] He, Z., Tran, K. P., Thomassey, S., Zeng, X., Xu, J., & Haiyi, C. (2020, December 29). A deep reinforcement learning based Multi-Criteria decision support system for textile manufacturing process optimization. arXiv.org. https://arxiv.org/abs/2012.14794

[5] Gwilt, A. (2020). A practical guide to sustainable fashion. https://doi.org/10.5040/9781350067059

[6] Liu, H., Li, C., Wu, Q., & Lee, Y. J. (2023, April 17). Visual instruction tuning. arXiv.org. https://arxiv.org/abs/2304.08485

[7] Burke, E. K., Gendreau, M., Hyde, M., Kendall, G., Ochoa, G., Özcan, E., & Qu, R. (2013). Hyper-heuristics: A survey of the state of the art. Journal of the Operational Research Society, 64(12), 1695–1724. https://doi.org/10.1057/jors.2013.71

[8] Jakobs, S. (1996). On genetic algorithms for the packing of polygons. European Journal of Operational Research, 88(1), 165–181. https://doi.org/10.1016/0377-2217(94)00166-9

[9] Kingma, D. P., & Welling, M. (2013, December 20). Auto-Encoding variational Bayes. arXiv.org. https://arxiv.org/abs/1312.6114

[10] Ruppert-Stroescu, M. (2025, January). Making Fashion Better: Rescue Waste and Reduce Invasive Species. International Textile and Apparel Association Annual Conference Proceedings. https://doi.org/10.31274/itaa.18871

[11] Aisyiyah, W. M., Soewardikoen, D. W., Azhar, H., & Nurhadiansyah, M. (2024, November). Innovation in patchwork waste processing using fabric manipulating techniques and design thinking approach. Advances in Social Humanities Research, 2(11), 1251–1258. https://doi.org/10.46799/adv.v2i11.307

[12] Jhariya, P. (2024, November). Integrating textile waste in interior design for sustainable furnishings: A case study of a cafe library. International Journal for Research in Applied Science and Engineering Technology, 12(11), 1025–1029. https://doi.org/10.22214/ijraset.2024.65273

[13] Song, D., Gao, D., Liu, G., & Li, X. (2024). FashionGPT: A large vision-language model for enhancing fashion understanding. Proceedings of the 33rd International Conference on Artificial Neural Networks (ICANN 2024), Lecture Notes in Computer Science, 15020, 308–323. Springer. https://doi.org/10.1007/978-3-031-72344-5_21

[14] Moreira, S., Felgueiras, H. P., & Marques, A. D. (2025, January). Circular Economy Practices in Fashion Design Education: The first phase of a case study. Sustainability, 17(3), 951. https://doi.org/10.3390/su17030951

[15] Ramsey, S., Grant, A., & Lee, J. (2025, January 26). Cross-Cultural fashion design via interactive large language models and diffusion models. arXiv.org. https://arxiv.org/abs/2501.15571

[16] Tu, X., Chen, D., Han, K., Altintas, O., & Wang, H. (2025, January 25). GreenAuto: An automated platform for sustainable AI model design on edge devices. arXiv.org. https://arxiv.org/abs/2501.14995

[17] Kim, J., et al. (2025, January). Inverse design of nanophotonic devices enabled by optimization algorithms and deep learning: Recent achievements and future prospects. Nanophotonics. https://doi.org/10.1515/nanoph-2024-0536

[18] Shao, G., et al. (2025, January). Reliable, efficient, and scalable photonic inverse design empowered by physics-inspired deep learning. Nanophotonics. https://doi.org/10.1515/nanoph-2024-0504

[19] Harjani, C. (2025, January). Zero-waste Pattern for Breastfeeding Mothers' Clothes According to the Style of Today's Working Women. KSS, 10(2), 97–103.

[20] Chen, G. (2024, December). Textile Product Innovative Design: The Practical Exploration of Integrating Fashion with Sustainable Development. AS, 3(6), 18–25.

[21] L. Yang, Y. Li, and L. Chen, "ClothPPO: A Proximal Policy Optimization Enhancing Framework for Robotic Cloth Manipulation with Observation-Aligned Action Spaces," *arXiv preprint arXiv:2405.04549*, May 5, 2024. [Online]. Available: https://arxiv.org/abs/2405.04549

[22] S. Bian, C. Xu, Y. Xiu, A. Grigorev, Z. Liu, C. Lu, M. J. Black, and Y. Feng, "ChatGarment: Garment estimation, generation and editing via large language models," *arXiv:2412.17811v3 [cs.CV]*, Apr. 2025. [Online]. Available: https://chatgarment.github.io/

[23] J. Yu, J. Li, Z. Yu, and Q. Huang, "Multimodal Transformer with Multi-View Visual Representation for Image Captioning," presented at the International Conference on Multimedia and Expo (ICME), Jul. 2015.

[24] Donati, L., Cesano, S., & Prati, A. (2019, July). A complete hand-drawn sketch vectorization framework. Multimedia Tools and Applications, 78. https://doi.org/10.1007/s11042-019-7311-3

[25] Sbai, O., Elhoseiny, M., Bordes, A., LeCun, Y., & Couprie, C. (2018, April 3). DeSIGN: Design Inspiration from Generative Networks. arXiv.org. https://arxiv.org/abs/1804.00921

[26] Chia, P. J., et al. (2022, April 8). Contrastive language and vision learning of general fashion concepts. arXiv.org. https://arxiv.org/abs/2204.03972